



MASTER OF SCIENCE OF INFORMATION SYSTEMS
INTELLIGENT SYSTEMS

Natural Language Processing Sentiment Analysis - AfriSenti Twitter Sentiment dataset

STUDENTS (GROUP2)	AINEDEMBE DENIS	2024-M132-23999
	MUSINGUZI BENSON	2024-M132-23947
LECTURER	Dr. Sibitenda Harriet	

About the AfriSenti Twitter Sentiment Dataset

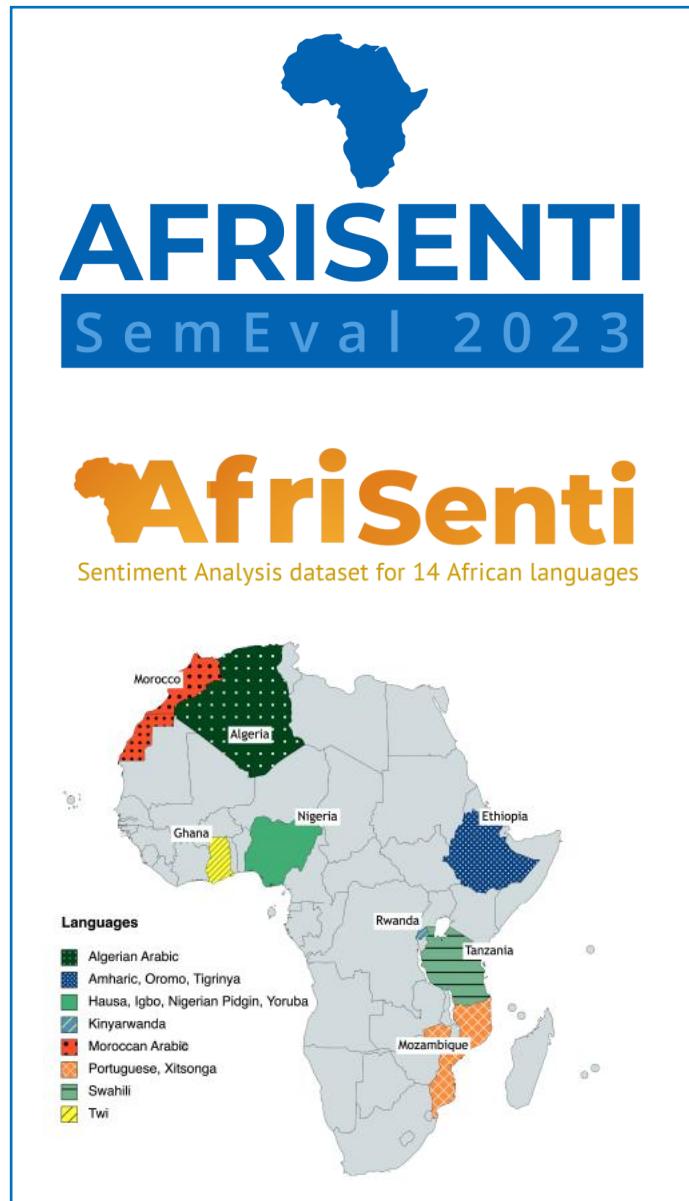
- The dataset contains 110,000+ annotated tweets across 14 African languages; Amharic, Algerian Arabic, Hausa, Igbo, Kinyarwanda, Mozambican Portuguese, Nigerian Pidgin, Oromo, Swahili, Tigrinya, Twi, Xitsonga, Yoruba
- Tweets are labeled by three annotators into three sentiment categories: Positive, Negative, Neutral
- All tweets are anonymized (@user) and URLs removed to protect user privacy
- Designed for Monolingual and multilingual and sentiment analysis, supporting research on African languages which are underrepresented in NLP.

Why this dataset matters:

- African languages make up 30% of the world's languages, yet lack NLP datasets (UNESCO) - AfriSenti fills this gap

References:

- Codalab: <https://codalab.lisn.upsaclay.fr/competitions/7320>,
- Hugging Face: <https://huggingface.co/datasets/HausaNLP/AfriSenti-Twitter>
- ACL Anthology: <https://aclanthology.org/2023.emnlp-main.862.pdf>



Project Objectives

Main Objective

- To build and evaluate a multilingual sentiment analysis model that classifies African-language tweets into positive, negative, or neutral sentiment.

Specific Objectives

- **Explore the dataset:** Analyze language distribution, tweet length, and sentiment label imbalance.
- **Preprocess the text:** Clean tweets, remove URLs, normalize slang, and tokenize using mBERT or XLM-RoBERTa.
- **Model Development:** Fine-tune XLM-RoBERTa for 3-class sentiment classification and compare with an LSTM baseline.
- **Model Training:** Use early stopping and gradient clipping for 3–5 epochs to prevent overfitting.
- **Evaluation:** Measure Accuracy, Macro-F1, ROC-AUC, and produce a confusion matrix. Include example predictions and attention visualization.
- **Ablation Studies:** Vary batch size, learning rate, and sequence length to observe performance changes.
- **Cross-Lingual Testing:** Train on one language (e.g., Swahili) - test on another (e.g., Amharic) to assess the model's multilingual transfer ability.

Real-Life Applicability

1. Social Media Monitoring

Governments, NGOs, and research institutions can track public mood on elections, health, crises, and social issues using multilingual sentiment trends.

2. Customer & Brand Analysis

Businesses can monitor consumer reactions across languages to improve products and services.

3. Hate Speech & Online Safety

Sentiment analysis is a key component in detecting negative or harmful content, especially in multilingual online spaces.

4. Policy & Public Opinion Research

Analysts can understand public attitudes in African countries/languages that are often ignored due to lack of datasets.

5. Enhancing African NLP Tools

AfriSenti supports the development of digital tools for African languages, promoting inclusive and equitable AI technology

6. Cross-Lingual AI Applications

Your cross-lingual experiments contribute to systems that work even when certain languages lack large datasets.

Loading AfriSenti-Twitter Dataset: Listing Available Language Configurations

LISTING AVAILABLE LANGUAGE CONFIGURATIONS FOR AfriSenti-Twitter DATASET

```
Available language configs: ['amh', 'hau', 'ibo', 'arq', 'ary', 'yor', 'por', 'twi', 'tso', 'tir', 'orm', 'pcm', 'kin', 'swa']
```

```
Loading Amharic (amh) with all splits
```

```
-----  
DatasetDict({  
    train: Dataset({  
        features: ['tweet', 'label'],  
        num_rows: 5984  
    })  
    validation: Dataset({  
        features: ['tweet', 'label'],  
        num_rows: 1497  
    })  
    test: Dataset({  
        features: ['tweet', 'label'],  
        num_rows: 1999  
    })  
})
```

```
Loading a single split /train only for Amharic (amh):
```

```
-----  
{'tweet': 'Tesfaye አብደ መ-በኩስ የተደራሱበትን ደቃ ለመፈከት እልማ የአሁን በቃ ነው ነው እናር ተንሽ', 'label': 2}
```

Language Code & Language Name Mapping:

amh	- Amharic
arq	- Algerian Arabic
ary	- Moroccan Arabic
hau	- Hausa
ibo	- Igbo
kin	- Kinyarwanda
orm	- Oromo
pcm	- Nigerian Pidgin
por	- Portuguese
swa	- Swahili
tir	- Tigrinya
tso	- Tsonga
twi	- Twi
yor	- Yoruba

Loading the AfriSenti Twitter Sentiment dataset

DATASET SAMPLE: One Example from Each Language

Loading one sample from each of 14 languages...

Lang.	Tweet	Sentiment
amh	Tesfaye ሌከስ ዘመን ለተደረሰኝን ይች ለጥቅና እወም የዚህ በቻ ነት ነ...	Positive
arq	@user على حسب موقعك يبدو أنك صاحب نظرة ثاقبة . يخي ...	Positive
ary	hhhhhhhhhhhhhhhhhh ana ga3ma sma3tt ach kant k...	Neutral
hau	@user Da kudin da Arewa babu wani abin azo agani d...	Positive
ibo	Nna Ike Gwuru ooo. 😊 https://t.co/NDS7juFBGd	Positive
kin	@user @user @user @user @user @user Hhhhhh n...	Positive
orm	@user Waa'ee mana waaqeffanaa ilaalcha keessa galc...	Neutral
pcm	yeah the guy wants to trend dat was why e join n...	Positive
por	Pedi uma resposta a Deus, ele deu me. Estou muito ...	Positive
swa	Kwani tanesco wanakataga umeme makusudinadhani kun...	Positive
tir	@user @user @user @user እንታይ ካብ ገዢ አዎረስምን ደሳ : በ...	Positive
tso	@user Loku u navela Ku tissunga, tissungue 🌐	Positive
twi	kako be shark but wo ti ewu	Positive
yor	Ìwọ ikú òpònú abaradúdú wọ, o ò şe é 're o. O d'ór...	Positive

Total languages shown: 14/14

Collect one example from each language

Label mapping: 0=negative,
1=neutral, 2=positive

Truncate tweet to 50 characters for better display

AfriSenti Twitter dataset: Total tweets across all languages

Language configs: ['amh', 'hau', 'ibo', 'arq', 'ary', 'yor', 'por', 'twi', 'tso', 'tir', 'orm', 'pcm', 'kin', 'swa']				
lang	train	validation	test	total
amh	5984	1497	1999	9480
hau	14172	2677	5303	22152
ibo	10192	1841	3682	15715
arq	1651	414	958	3023
ary	5583	494	2961	9038
yor	8522	2090	4515	15127
por	3063	767	3662	7492
twi	3481	388	949	4818
tso	804	203	254	1261
tir	0	398	2000	2398
orm	0	396	2096	2492
pcm	5121	1281	4154	10556
kin	3302	827	1026	5155
swa	1810	453	748	3011

Total tweets across all languages: 111718

Initial Data Exploration: Load train data from all languages

Dataset Loading Summary

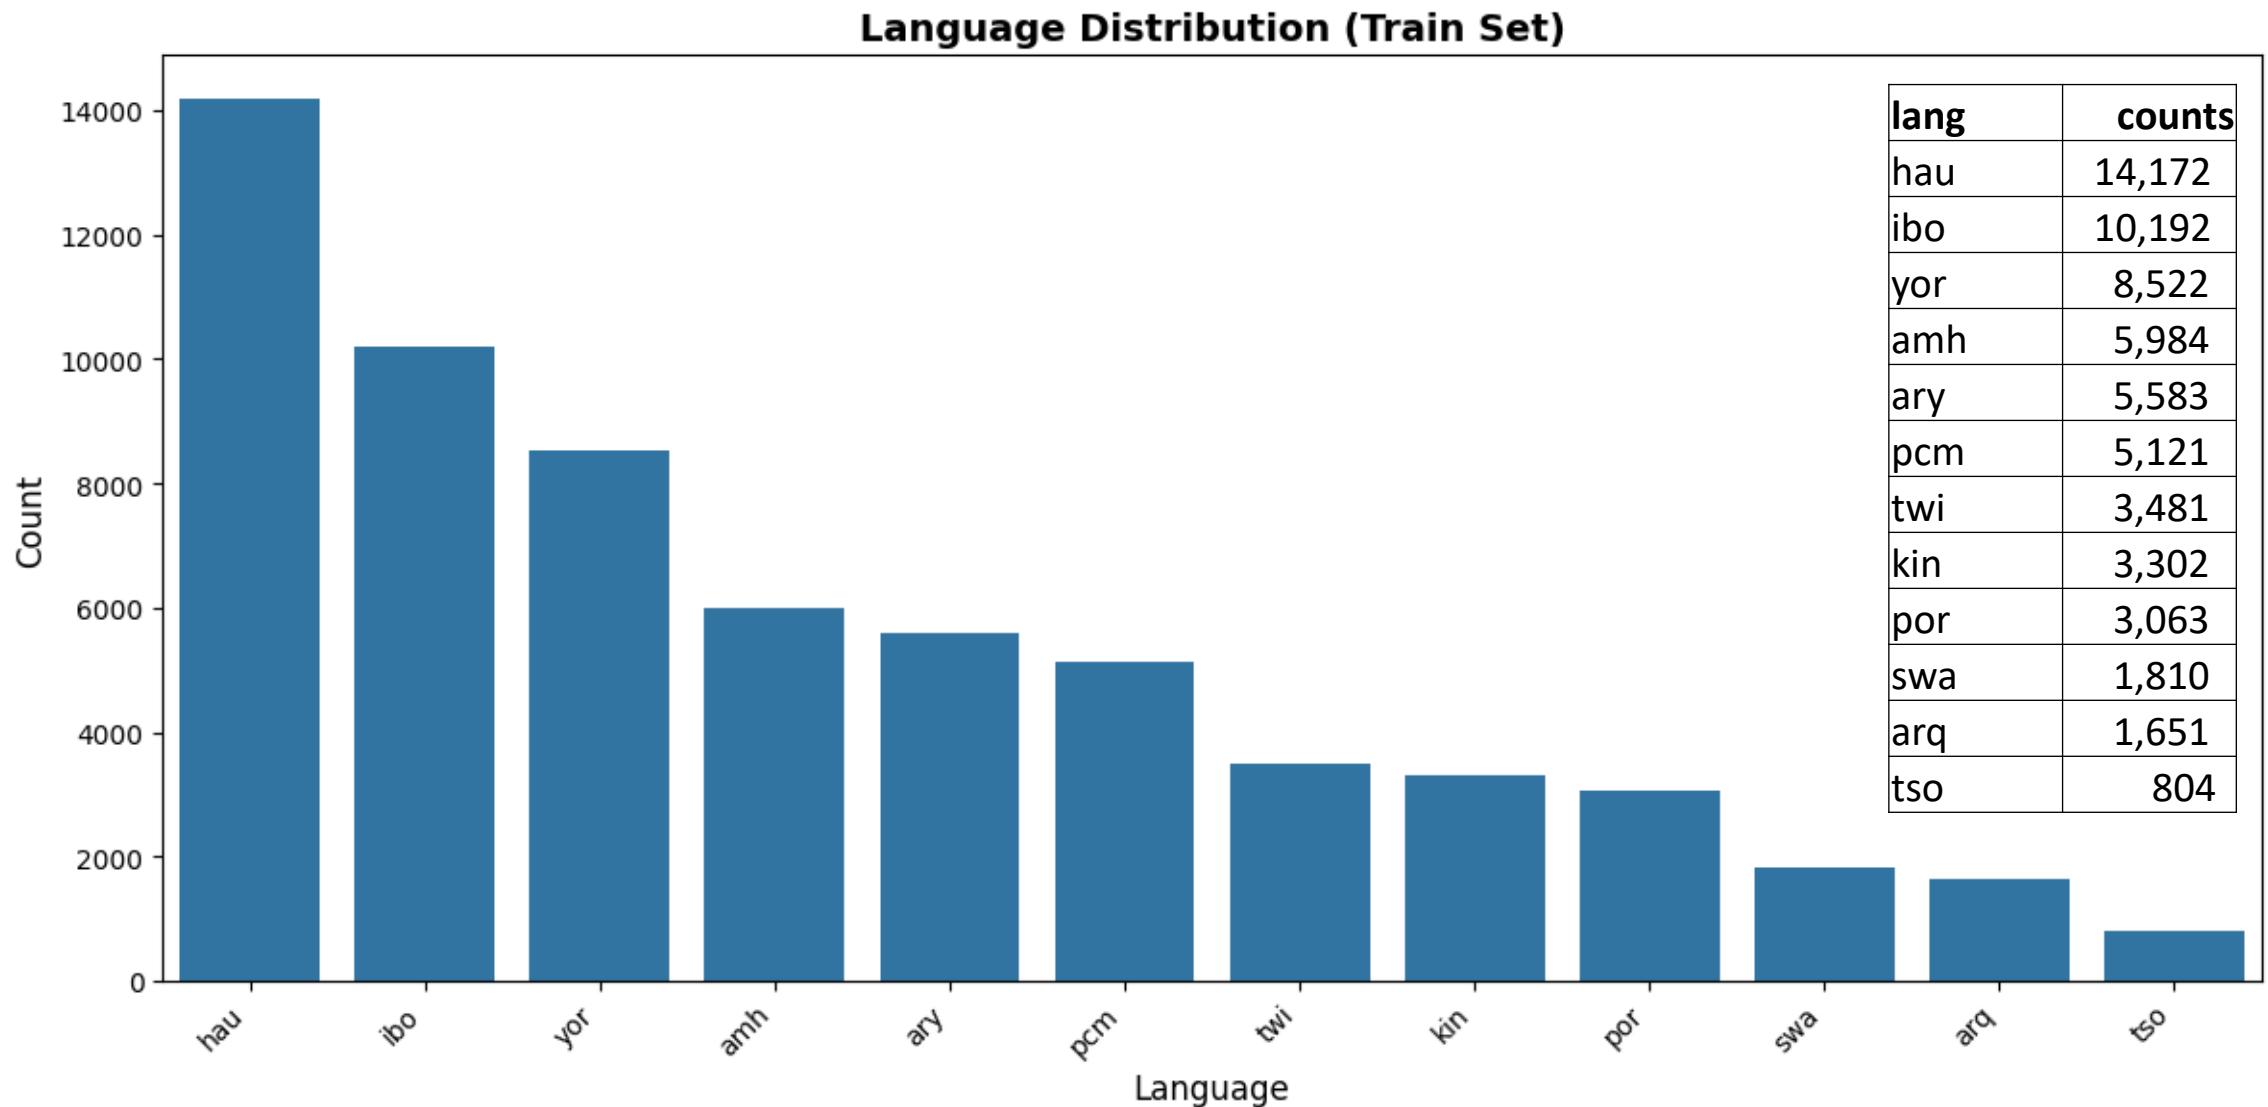
Language Code	Language Name	Train Samples	Status
amh	Amharic	5984	Loaded
hau	Hausa	14172	Loaded
ibo	Igbo	10192	Loaded
arq	Algerian Arabic	1651	Loaded
ary	Moroccan Arabic	5583	Loaded
yor	Yoruba	8522	Loaded
por	Portuguese	3063	Loaded
twi	Twi	3481	Loaded
tso	Tsonga	804	Loaded
tir	Tigrinya	0	No train split
orm	Oromo	0	No train split
pcm	Nigerian Pidgin	5121	Loaded
kin	Kinyarwanda	3302	Loaded
swa	Swahili	1810	Loaded

Total train samples loaded: 63,685

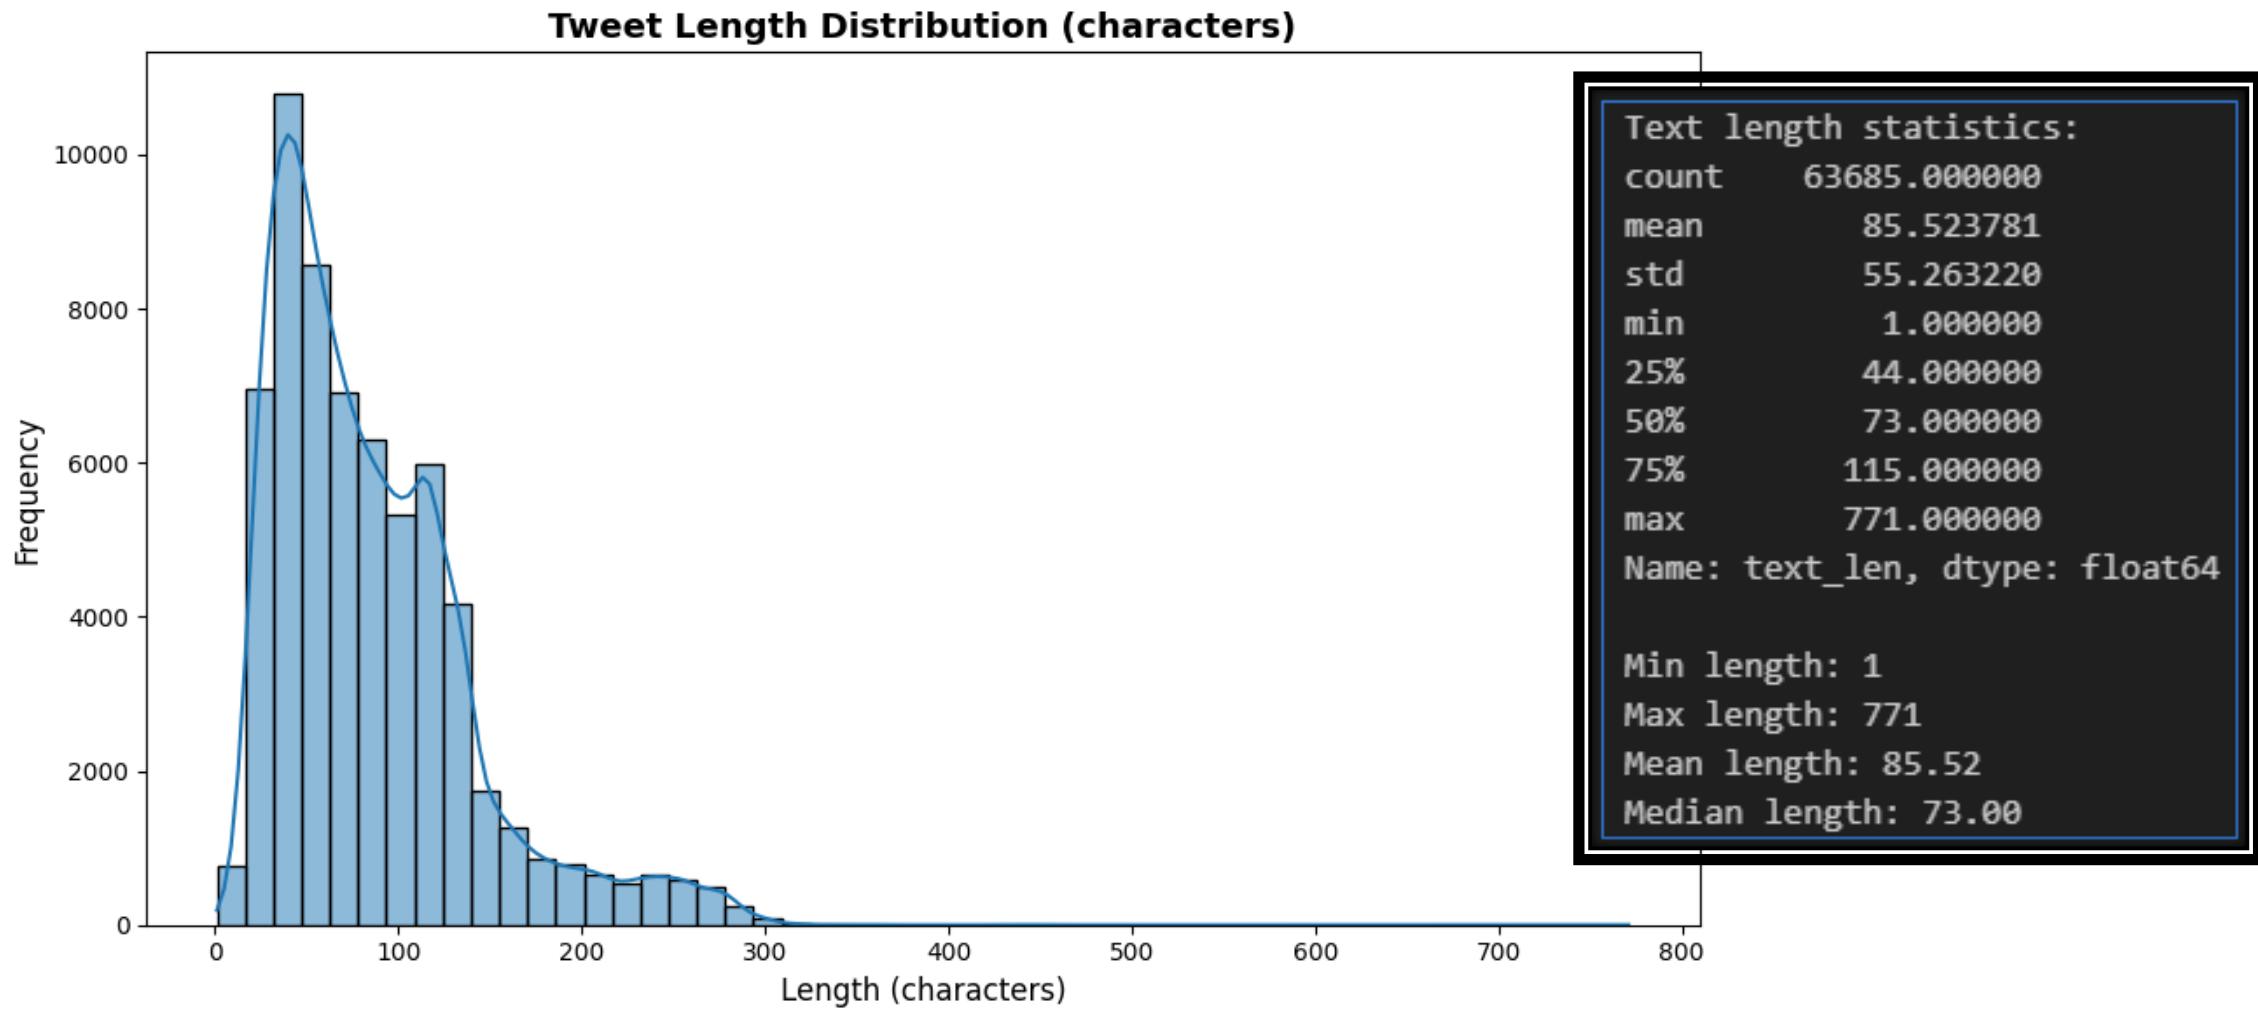
First few rows:

	lang	lang_name	tweet	label	label_text
0	amh	Amharic	Tesfaye ሊከ ቅዱስ ለበሰኑ የተደራሰሩን ደቶ ለጥሩ እልም ያለሁ በ...	2	positive
1	amh	Amharic	ይሂዱ ነው እናይል የእውቀትኩ ጥገ....በሰኑ ስጻ ከምትኝነት ለምን ታረክ...	2	positive
2	amh	Amharic	ዘግበ ይበላል? ለለ የሚባል ንር ካለ እንተው ገንዘብ!	2	positive
3	amh	Amharic	? ይደ በዘመን ከዚህ ደቶ በት ፍላሽ ቁ ሲሉ እየተችን +መኖች እንዲያው...	2	positive
4	amh	Amharic	የልማት?? ??? ገንዘብ	2	positive

Initial Data Exploration: Language Distribution on the Train Set

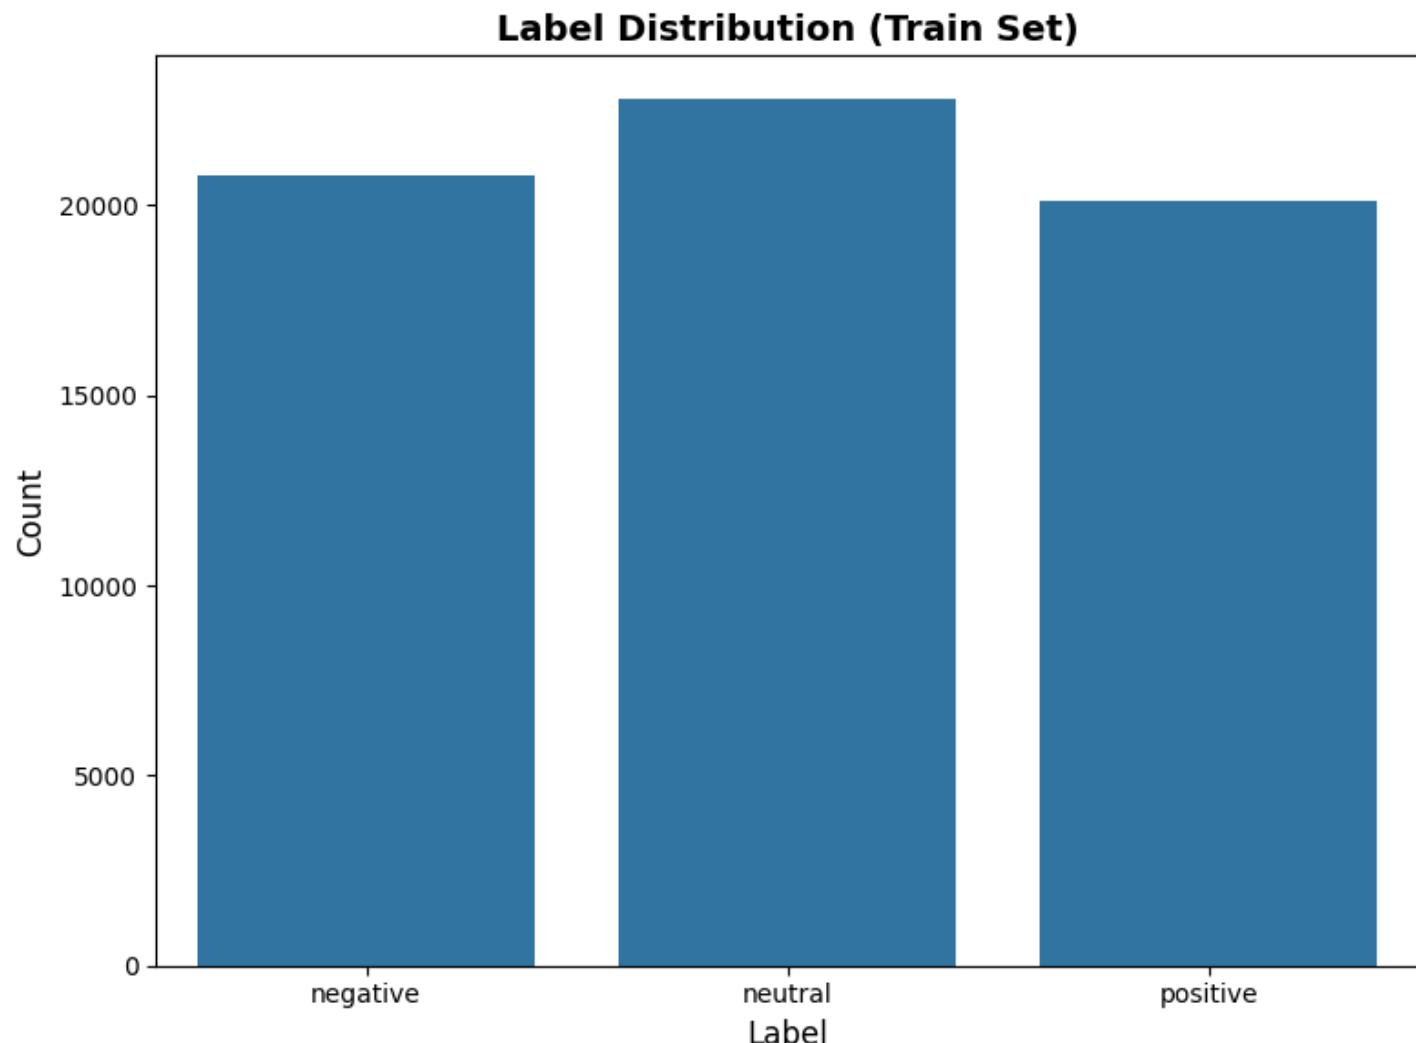


Initial Data Exploration: Text length Distribution the Train Set



The training tweets are mostly short (median 73 characters), with a right-skewed distribution and a few very long outliers. The data is concise and typical of Twitter usage, making it suitable for transformer models with moderate sequence lengths

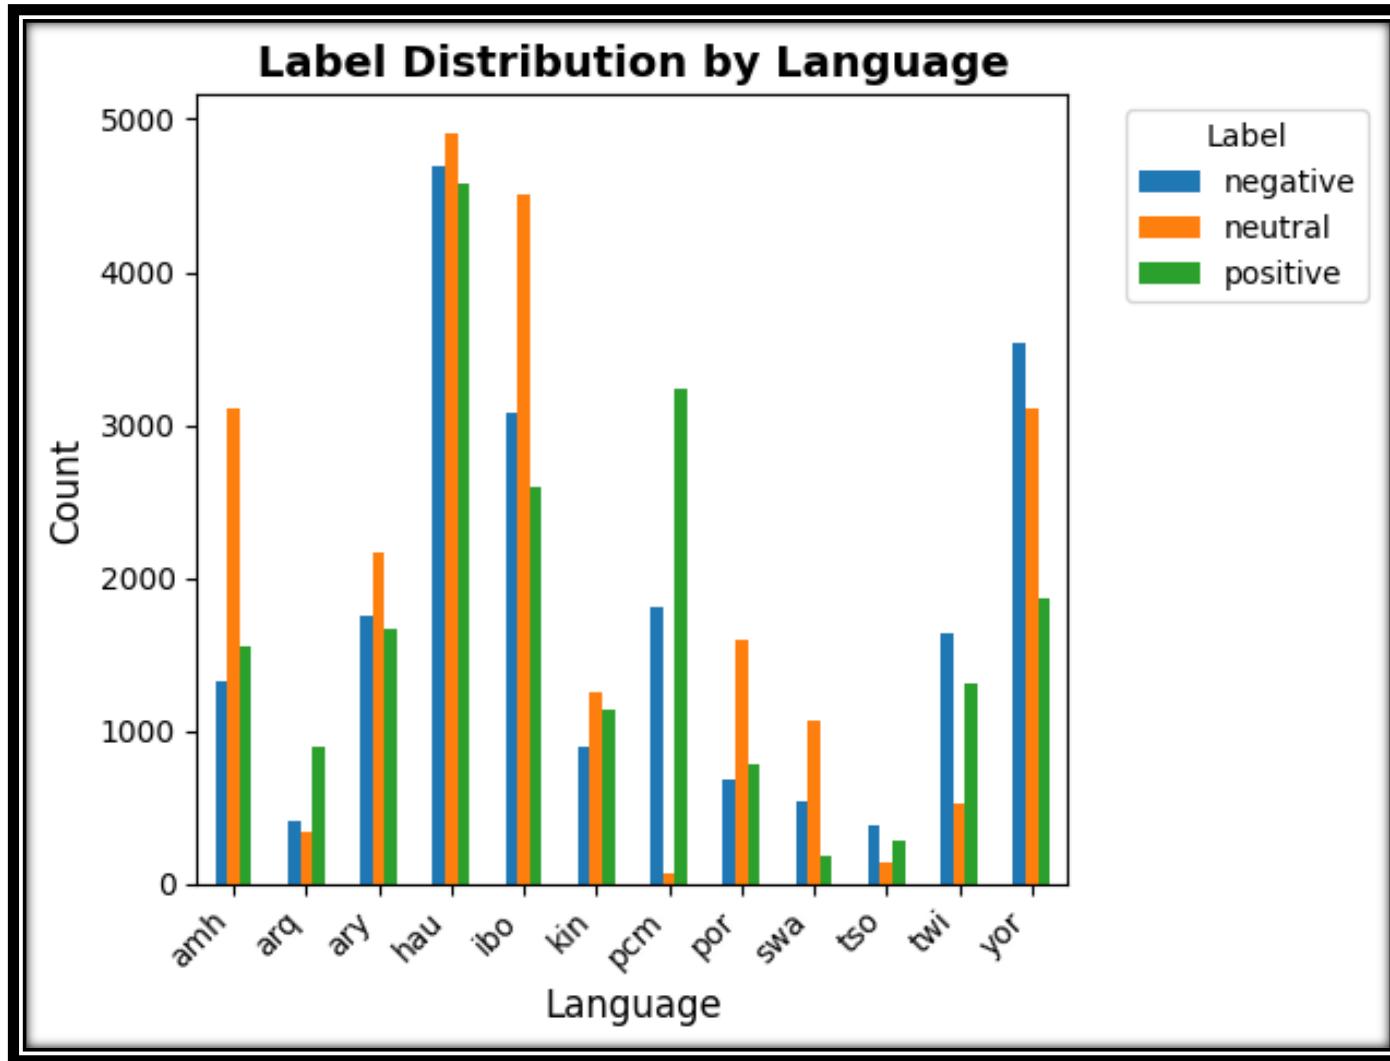
Initial Data Exploration: Overall Label Distribution



```
Label counts:  
label_text  
neutral      22794  
negative     20783  
positive     20108  
Name: count, dtype: int64  
  
Label percentages:  
label_text  
neutral      35.791788  
negative     32.634058  
positive     31.574154  
Name: proportion, dtype: float64
```

The train set shows a well-balanced label distribution with 36% neutral, 33% negative, and 32% positive tweets. This observed balance supports stable model training without the need for class reweighting.

Initial Data Exploration: Label Distribution by Language



Label distribution by language:			
label_text	negative	neutral	positive
lang			
amh	1332	3104	1548
arq	417	342	892
ary	1758	2161	1664
hau	4687	4912	4573
ibo	3084	4508	2600
kin	899	1257	1146
pcm	1808	72	3241
por	681	1600	782
swa	547	1072	191
tso	384	136	284
twi	1644	522	1315
yor	3542	3108	1872

- 1) Sentiment distribution varies significantly across languages
- 2) Some languages like (Hausa, Igbo, and Yoruba) are balanced, others (e.g., Nigerian Pidgin and Twi) show extreme skewness in one sentiment class.
- 3) Potential challenges for model training and emphasizes the importance of multilingual transfer and careful evaluation

Preprocess text using multilingual tokenizers (mBERT, XLM-RoBERTa)

Load train, validation, and test datasets from all languages

Train samples: 63685

Validation samples: 13726

Test samples: 34307

Text preprocessing. Handled emojis, URLs, and slang normalization.

Text preprocessing completed

Sample preprocessed tweets:

1. Tesfaye አከላ ማስፈልጊ የተደረሰበን ይች ለመፈጸም የዚህ በቻ ነው እናለን ተንስ
2. የሆነ ነው አይደል የእውቀትኩ ጥሩ....በዚህ ስሜ ከምትኩንኑ ለማን ታሪክ አጥነበዋም....ዚህ ሲስጡን አያዝነዋል
3. አገበ ይበላል? ለለ የሚበላል ነገር ካለ አንተው ገንዘን!

- URLs
- Mentions <USER>
- Normalize whitespace
- Remove # but keep word

Preprocess text: Label mapping & tokenizing Tweets Using XLM-RoBERTa

```
Label mappings:
```

- label2id: {'negative': 0, 'neutral': 1, 'positive': 2}
- id2label: {0: 'negative', 1: 'neutral', 2: 'positive'}

```
Map: 100%|██████████| 63685/63685 [00:00<00:00, 288476.51 examples/s]
Map: 100%|██████████| 13726/13726 [00:00<00:00, 108999.68 examples/s]
Map: 100%|██████████| 34307/34307 [00:00<00:00, 240070.99 examples/s]
```

```
Label encoding completed
```

Label mapping:

string -> int

(0=negative, 1=neutral, 2=positive)

Tokenizing Tweets

Using multilingual tokenizer
XLM-RoBERTa

```
Loading multilingual tokenizer: xlm-roberta-base
Tokenizer loaded successfully: xlm-roberta-base
Vocabulary size: 250,002
```

Max sequence length: 128

Tokenizing datasets with multilingual tokenizer

```
Map: 100%|██████████| 63685/63685 [00:07<00:00, 9038.61 examples/s]
Map: 100%|██████████| 13726/13726 [00:01<00:00, 8075.26 examples/s]
Map: 100%|██████████| 34307/34307 [00:03<00:00, 8613.86 examples/s]
```

Tokenization completed using multilingual tokenizer

Modeling

**Fine-tune XLM-RoBERTa for 3-class sentiment classification
(positive, neutral, negative).**

Compare with

LSTM baseline Model

XLM-RoBERTa Transformer Model

```
Device set to: cpu
Loading XLM-RoBERTa model: xlm-roberta-base

- Number of labels: 3
- Label mappings: {'negative': 0, 'neutral': 1, 'positive': 2}

Some weights of XLMRobertaForSequenceClassification were not initialized
You should probably TRAIN this model on a down-stream task to be
able to use it.

Successfully loaded the xlm-roberta-base model
Total parameters: 278,045,955
```

Total parameters:
278,045,955

LSTM Baseline Model

LSTM Model Configuration:

- Vocabulary size: 250,002
- Embedding dimension: 128
- Hidden dimension: 128
- Number of labels: 3

LSTM model created

- Embedding dim: 128
- Hidden dim: 128
- Total parameters: 32,265,219

Total parameters:
32,265,219

Comparison: XLM-RoBERTa vs LSTM Baseline

PARAMETER COMPARISON:

- XLM-RoBERTa has 8.6x more parameters than LSTM Baseline
- XLM-RoBERTa: 278,045,955 parameters (pre-trained, fine-tuned on AfriSenti)
- LSTM Baseline: 32,265,219 parameters (trained from scratch on AfriSenti)
- Difference: 245,780,736 parameters (761.8% more)

Note: XLM-RoBERTa's larger parameter count reflects its pre-trained multilingual knowledge, while LSTM is a lighter baseline model trained only on this dataset.

Training Models - 5 epochs with Early stopping + Gradient clipping)

Training LSTM Baseline Model with early stopping

Training LSTM Baseline Model

```
[LSTM] Epoch 1/3 | Train loss: 0.9572 | Val acc: 0.5762 | Val F1: 0.5683  
- New best F1: 0.5683, model saved!
```

```
[LSTM] Epoch 2/3 | Train loss: 0.7799 | Val acc: 0.6189 | Val F1: 0.6160  
- New best F1: 0.6160, model saved!
```

```
[LSTM] Epoch 3/3 | Train loss: 0.6447 | Val acc: 0.6282 | Val F1: 0.6262  
- New best F1: 0.6262, model saved!
```

```
LSTM training completed. Best F1: 0.6262
```

Training Models - 5 epochs with Early stopping + Gradient clipping)

Training **XLM-RoBERTa** or AfriBERTa Model with early stopping

```
Training Transformer Model (xlm-roberta-base)
```

```
-----
```

```
[TRANS] Epoch 1/3
```

```
    Train loss: 0.8997 | Val acc: 0.6241 | Val F1: 0.6242
```

```
    New best F1: 0.6242, model saved!
```

```
[TRANS] Epoch 2/3
```

```
    Train loss: 0.7241 | Val acc: 0.6689 | Val F1: 0.6688
```

```
    New best F1: 0.6688, model saved!
```

```
[TRANS] Epoch 3/3
```

```
    Train loss: 0.6122 | Val acc: 0.6792 | Val F1: 0.6795
```

```
    New best F1: 0.6795, model saved!
```

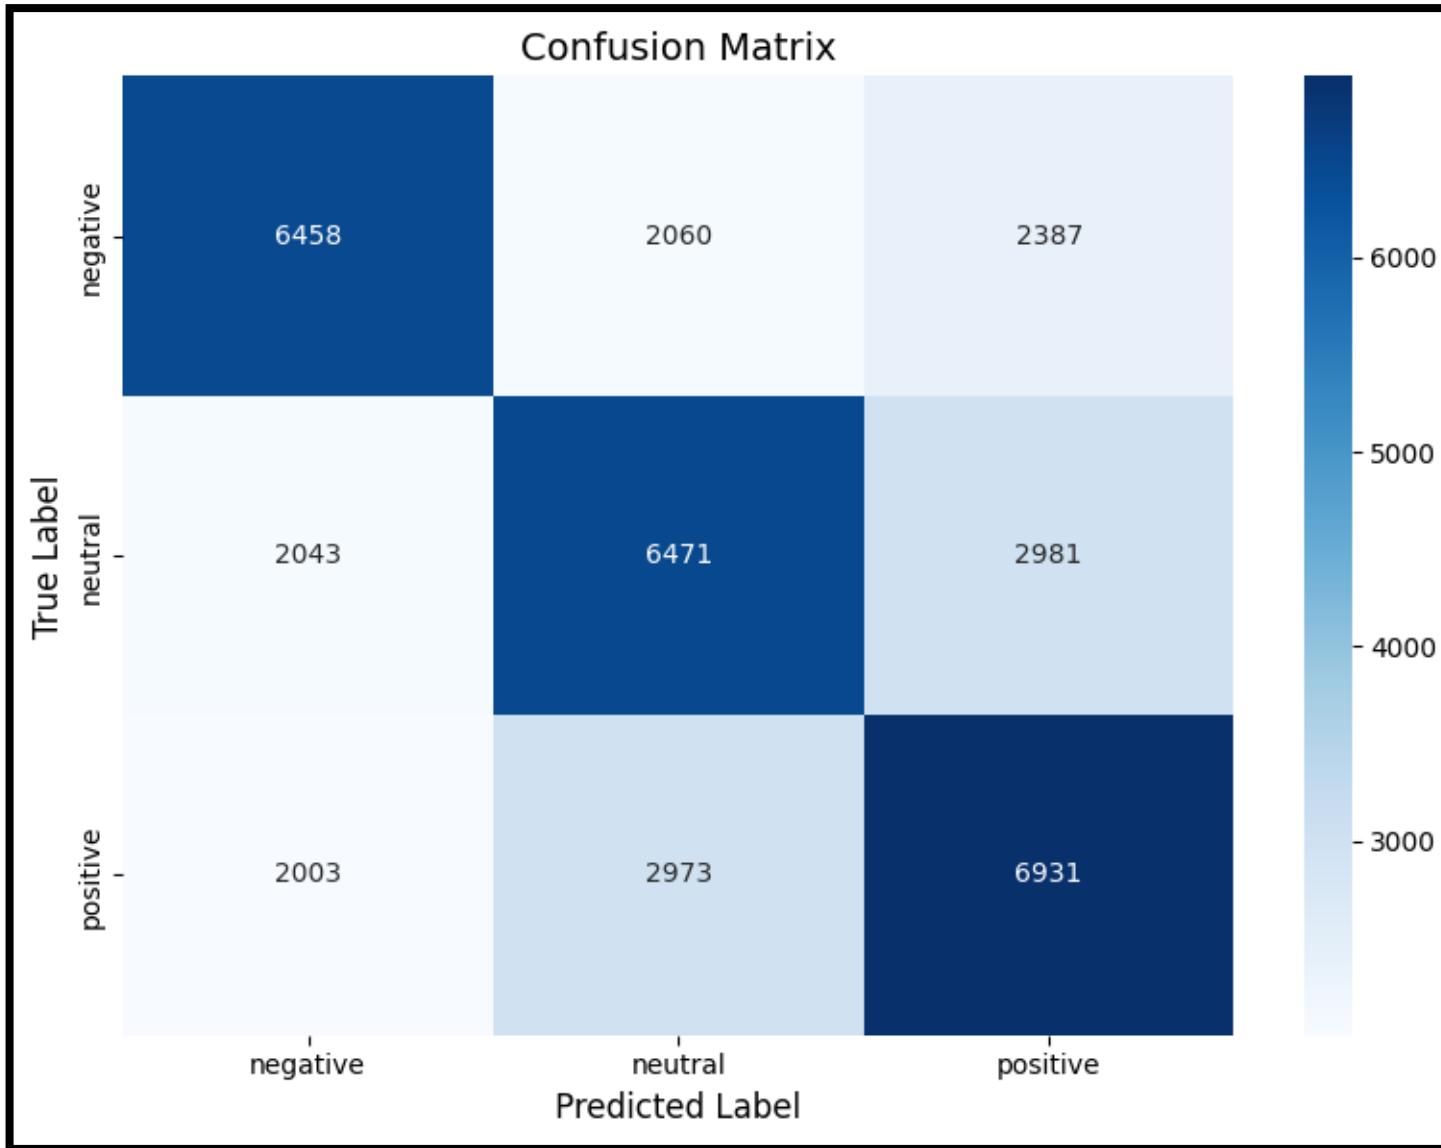
```
Transformer training completed. Best F1: 0.6795
```

```
Total training time: 91.3 minutes (1.52 hours)
```

Evaluation (F1, Accuracy, ROC-AUC, Confusion Matrix) + Predictions & Attention

Evaluating both LSTM and XLM-RoBERTa models on test data with comprehensive metrics, example predictions, and attention visualization.

Evaluating LSTM Baseline Model on Test Set

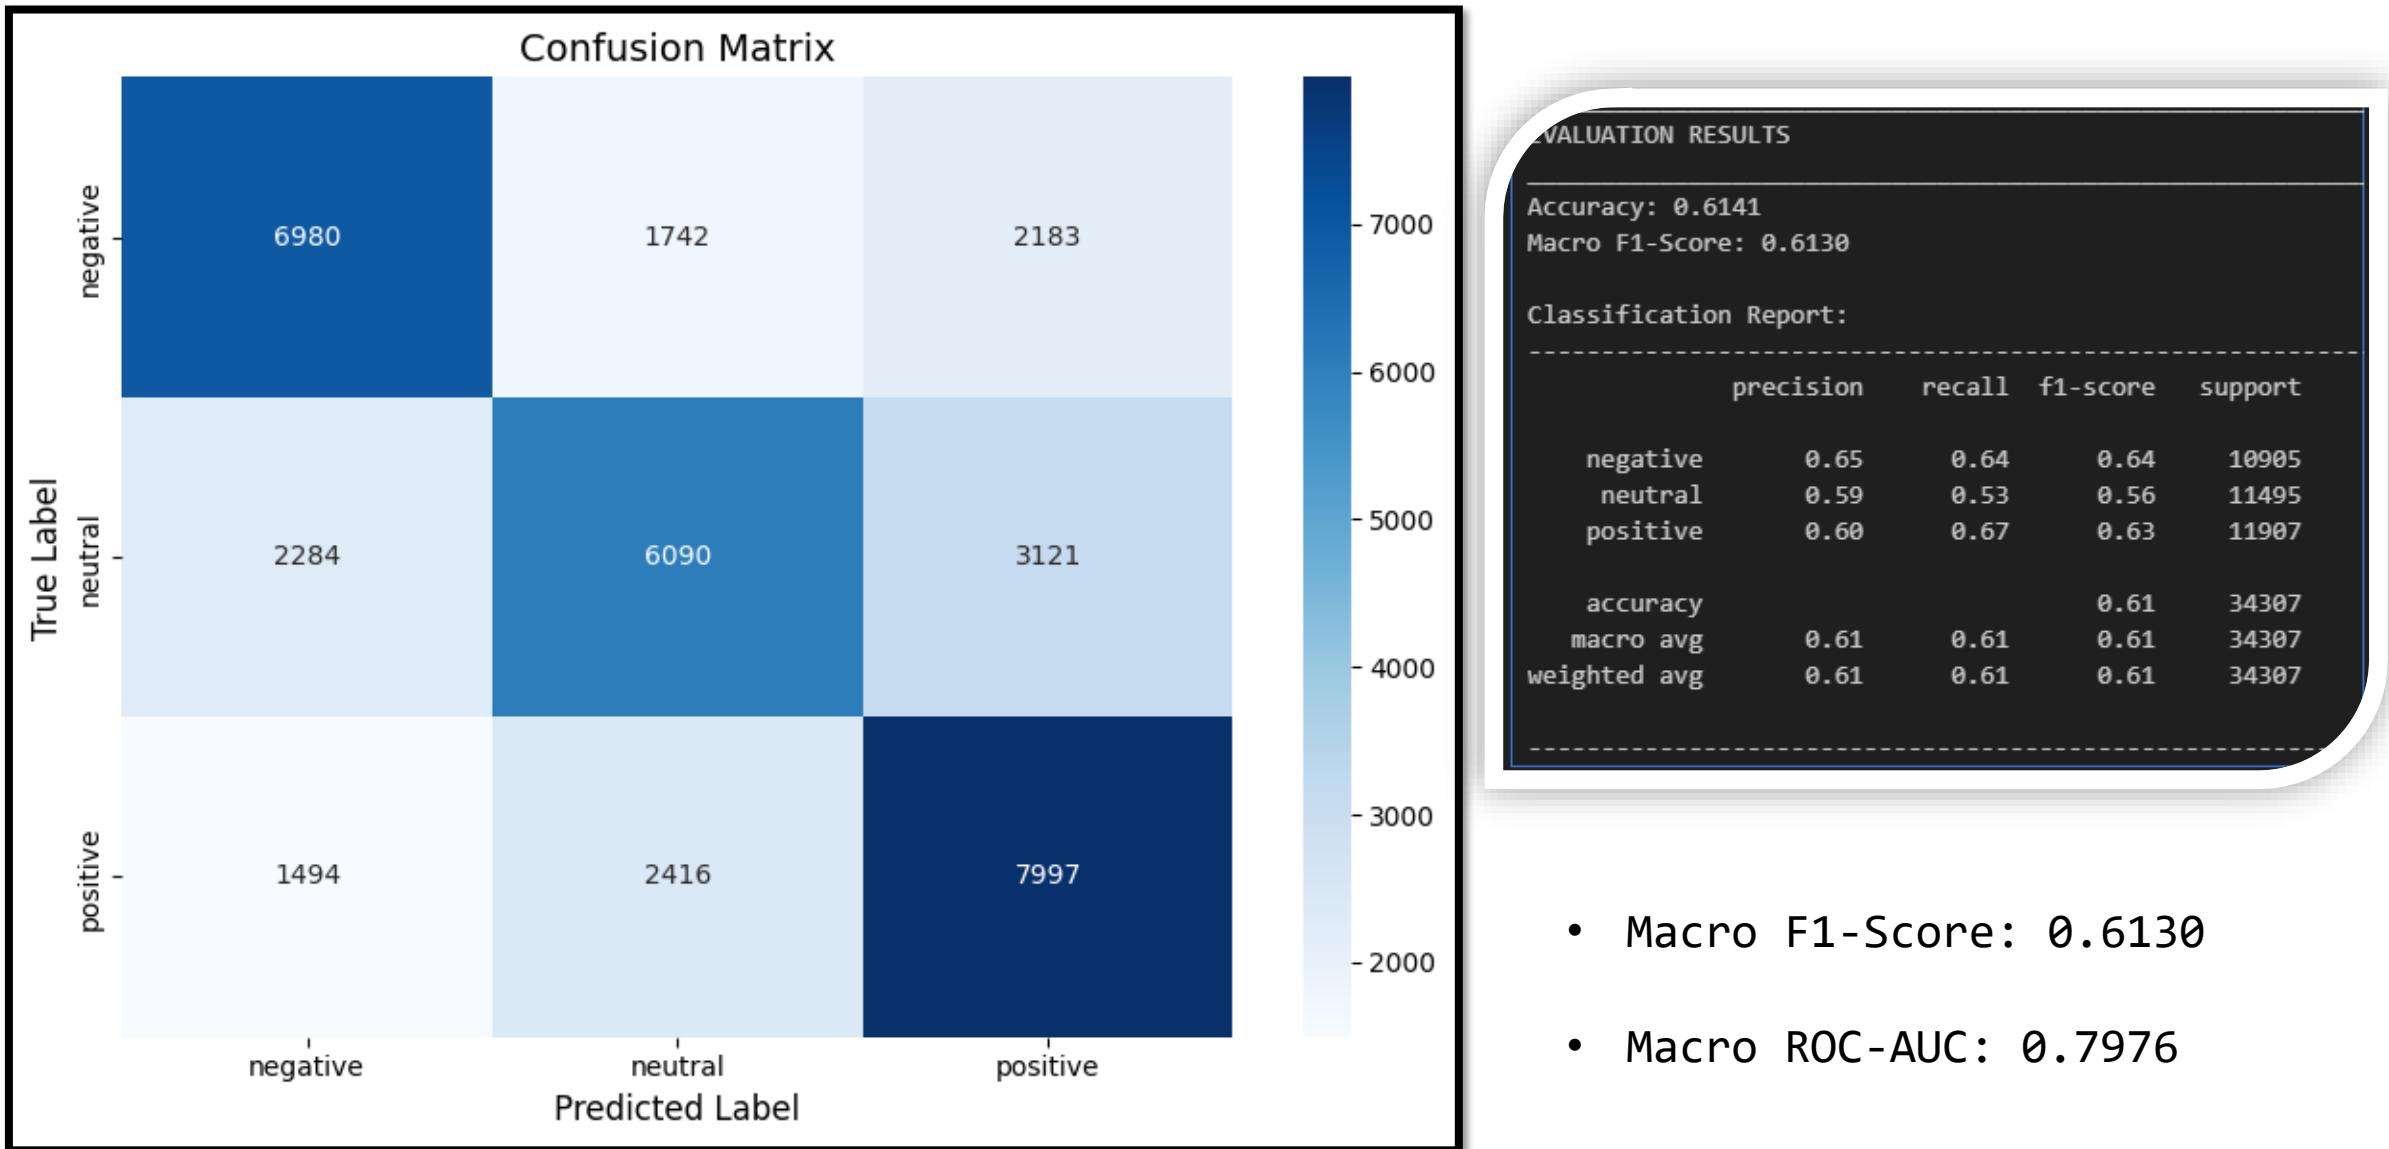


EVALUATION RESULTS

Accuracy:	0.5789			
Macro F1-Score:	0.5796			
Classification Report:				
	precision	recall	f1-score	support
negative	0.61	0.59	0.60	10905
neutral	0.56	0.56	0.56	11495
positive	0.56	0.58	0.57	11907
accuracy			0.58	34307
macro avg	0.58	0.58	0.58	34307
weighted avg	0.58	0.58	0.58	34307

- Macro F1-Score: 0.5796
- Macro ROC-AUC: 0.7590

Evaluating XLM-ROBERTA Transformer Model on Test Set



Predictions on Sample Texts do demonstrate model Behavior on Different Languages and Sentiment Classes

EXAMPLE PREDICTIONS - XLM-ROBERTA TRANSFORMER MODEL

Text: Nimefurahi sana kwa huduma hii! 😊

Predicted Sentiment: negative

Probabilities:

- Negative: 0.9881
- Neutral: 0.0101
- Positive: 0.0017

Text: Service hii ni mbaya sana.

Predicted Sentiment: positive

Probabilities:

- Negative: 0.0122
- Neutral: 0.0201
- Positive: 0.9677

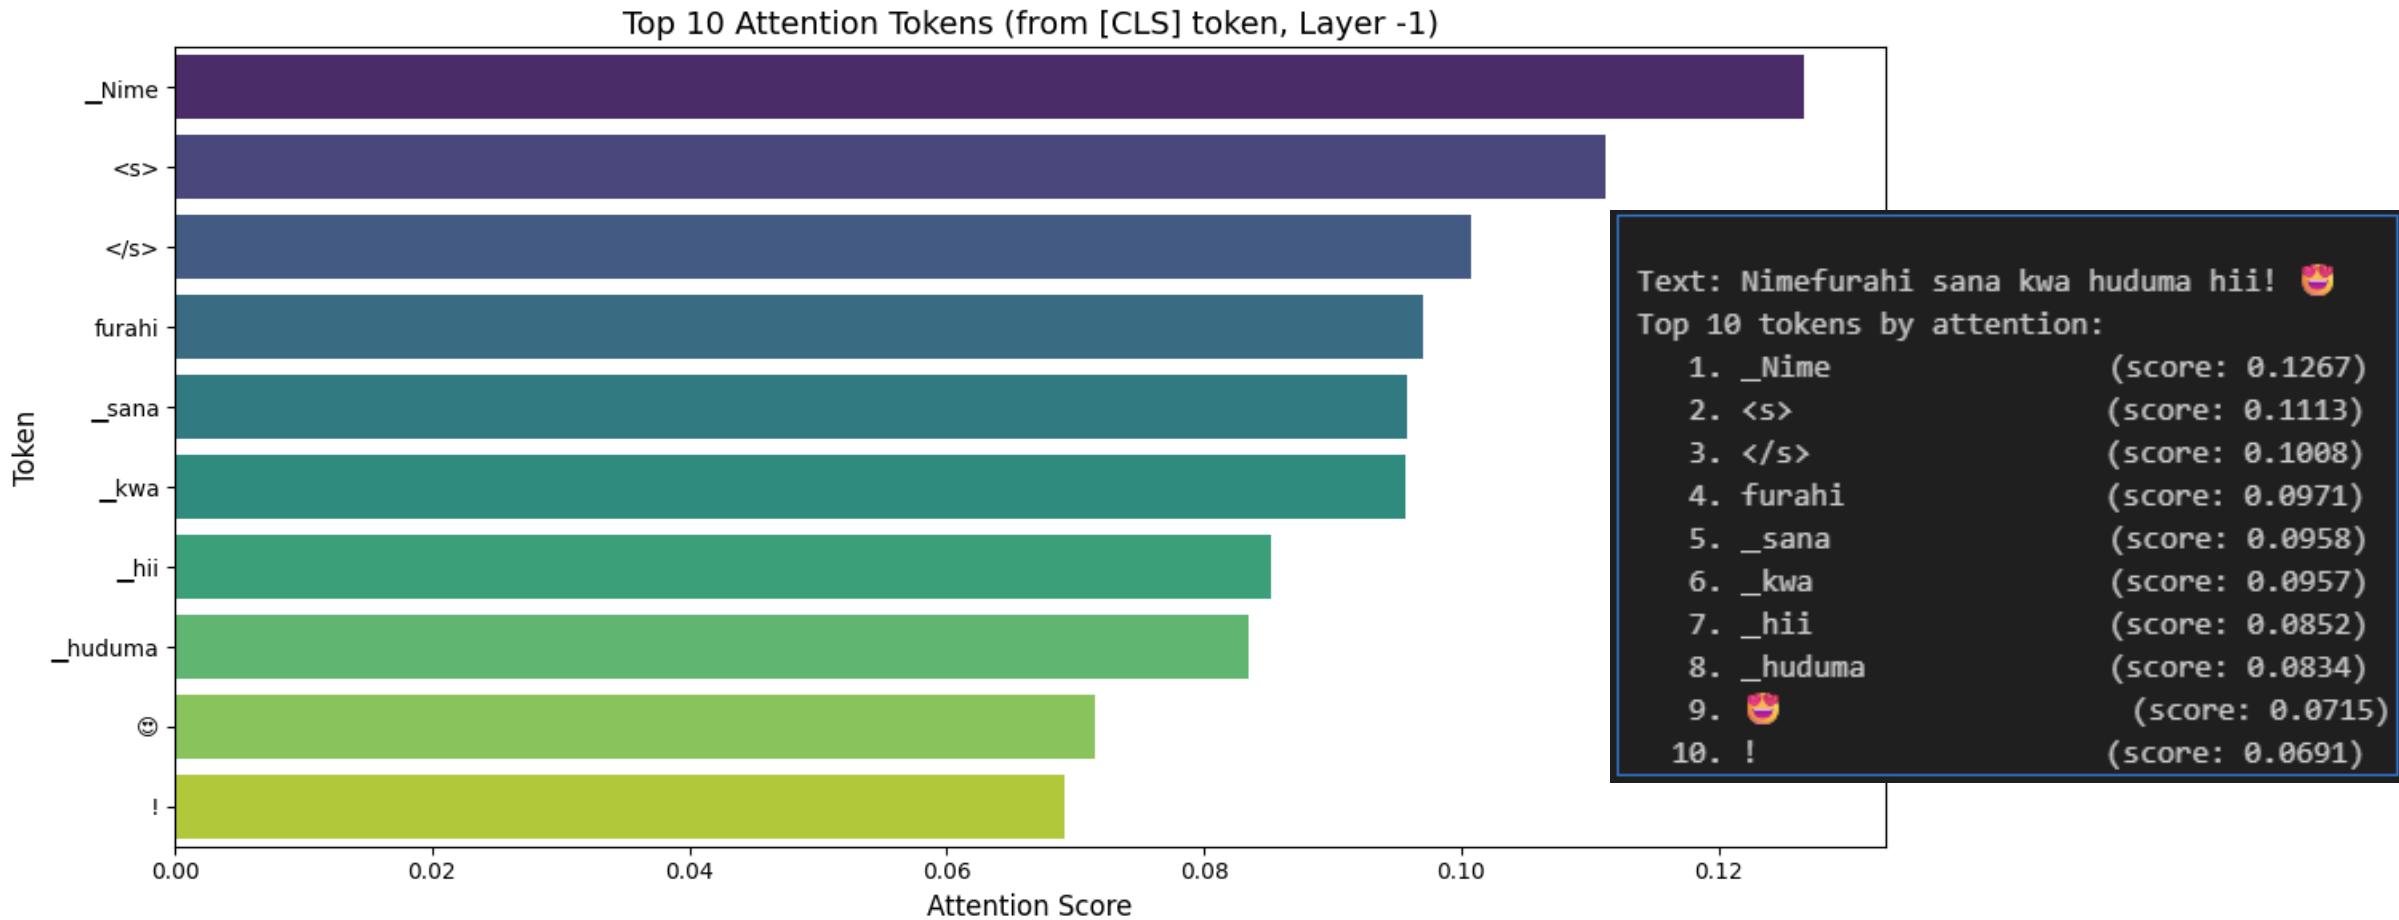
Text: I am not sure how I feel about this.

Predicted Sentiment: positive

Probabilities:

- Negative: 0.0107
- Neutral: 0.0747
- Positive: 0.9146

Attention Visualization XLM-ROBERTA Transformer Model



Ablation Studies (Batch Size, Learning Rate, Sequence Length)

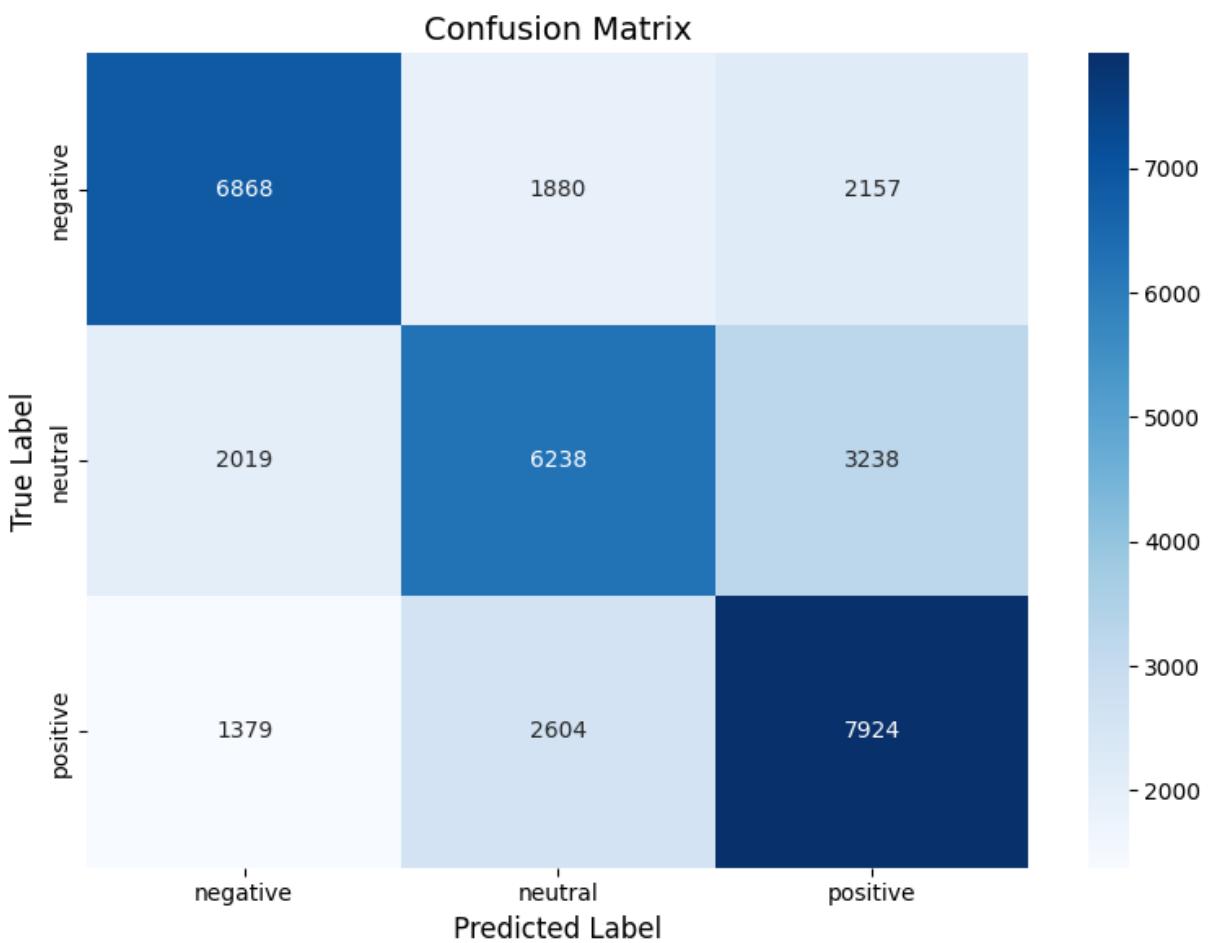
```
Training model with config: batch_size=8, lr=2e-05, max_length=128
-----
[TRANS] Epoch 1/3
  Train loss: 0.8998 | Val acc: 0.6357 | Val F1: 0.6310
  New best F1: 0.6310, model saved!
[TRANS] Epoch 2/3
  Train loss: 0.7203 | Val acc: 0.6628 | Val F1: 0.6606
  New best F1: 0.6606, model saved!
[TRANS] Epoch 3/3
  Train loss: 0.5979 | Val acc: 0.6783 | Val F1: 0.6789
  New best F1: 0.6789, model saved!
Transformer training completed. Best F1: 0.6789
Total training time: 117.6 minutes (1.96 hours)
EVALUATION RESULTS

Accuracy: 0.6130
Macro F1-Score: 0.6129

Classification Report:
-----
             precision    recall  f1-score   support

  negative       0.67      0.63      0.65     10905
  neutral        0.58      0.54      0.56     11495
  positive        0.59      0.67      0.63     11907

  accuracy           -         -      0.61    34307
  macro avg        0.62      0.61      0.61    34307
  weighted avg     0.61      0.61      0.61    34307
```



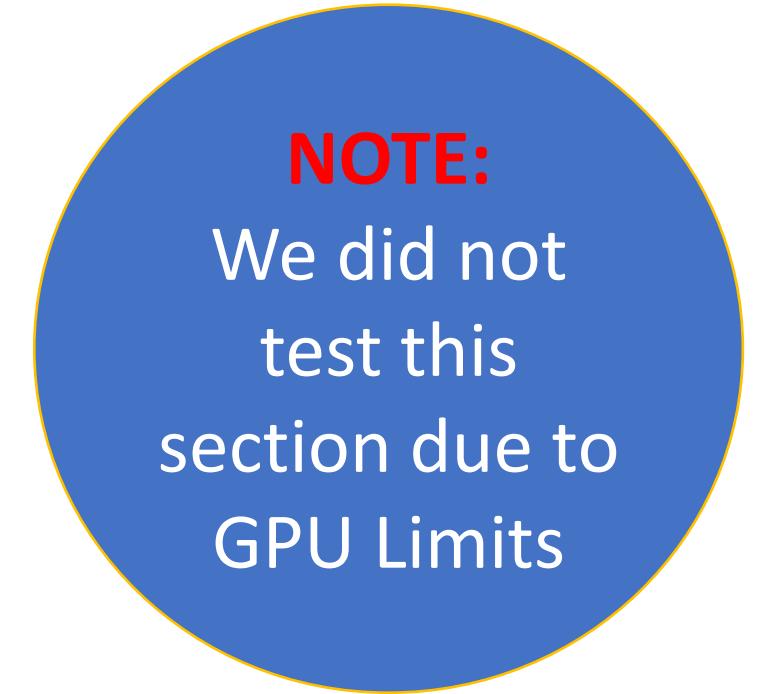
NOTE: This study did not complete due to GPU Limits

Cross-Lingual Testing: Multiple Language Pairs

Test cross-lingual transfer with multiple language pairs:

- Train on Swahili, Test on Amharic
- Train on Swahili, Test on Pidgin English (pcm)

Testing the model's ability to transfer knowledge across languages by training on one language (Swahili) and evaluating on different target languages.



THE END

STUDENTS	AINEDEMBE DENIS +256 788-674576 dembedenisjb@gmail.com, ainedembe.denis@stud.umu.ac.ug
	MUSINGUZI BENSON +256 782 942245, musiben@gmail.com, musinguzi.benson@stud.umu.ac.ug
LECTURER	Dr. Sibitenda Harriet +256 777 056581 hsibitenda@umu.ac.ug